

check 3

2. The TCP Sender

TCP:一种网络协议，它在不可靠的网络数据流上实现一对流控制的可靠传输的字节流。每一层的对等体经行流控制。

TCPSender:

- 跟踪接收方窗口大小
- 读数据，创建新的TCP数据段，尽可能填满接收方窗口，发送数据。
- 跟踪已发送但未被确认的数据段
- 重发经过既定时间仍未被确认的数据段

2.1 How does the TCPSender know if a segment was lost?

- TCPSender 的tick函数计算时间的流逝，超时便重传。
- RTO使用`_initial_retransmission_timeout`进行初始化，RTO会随时间变化
- 实现超时重传计时器timer，只能使用tick
- 每次发送数据都要启动计时器，RTO事件后过期。
- 已调用tick并且计时器过期：
 - 重传最早的未被确认的段；这些数据要保存在特定的数据结构中
 - 如果窗口大小非零：
 - 跟踪连续重传的数据并增加它，TCPConnection要使用此信息判断是否连接无望并终止
 - RTO值加倍（指数回退），减缓重传速度
 - 重置计时器，在RTO时间后过期（考虑RTO的变化）
- 当接收方给发送方一个确认成功接收到新数据的确认(该确认反映了一个比之前任何确认都大的绝对序列号):
 - 将RTO设为初始值
 - 发送方有未完成的数据，重启计时器
 - 将连续重传计数设为0

2.2 Implementing the TCP sender

TCPSender的基本思想：

- 给定的数据字节流
- 分割成段
- 发送给接收方
- 没有很快得到确认，重新发送未被确认的数据

具体接口：

```
1 void push( Reader& outbound_stream );
```

从流中读取数据，只要有新的字节流过来并且有足够的空间，就尽可能多的生成TCPSenderMessage，每个TCPSenderMessage不能大于TCPConfig::MAX PAYLOAD SIZE（1452字节），TCPSenderMessage::sequence length() 用来计算一个TCPSenderMessage所占的序列号数量；如果接收方宣布接收窗口为0，发送方则需要假设窗口大小为1，发送数据，可能会受到接收方反馈的最新窗口大小，继续进行数据传输。

```
1 std::optional<TCPSenderMessage> maybe send();
```

TCPSender实际发送TCPSenderMessage的机会

```
1 void receive( const TCPReceiverMessage& msg );
```

从接收方接受消息，传递新的窗口边界。查看未完成的数据段并清楚已完成的数据段（确认过）。

```
1 void tick( const size_t ms_since_last_tick );
```

记录时间流逝情况。

```
1 void send_empty_message();
```

发送空消息，接收方要配合

▼ tcp_sender.hh

```
1 class TCPSender
2 {
3     Wrap32 isn_;
4     bool is_send_ISN; /*the flag of ISN*/
5     bool is_send_FIN; /*the flag of FIN*/
6     uint64_t initial_RTO_ms_;
7     int cur_RTO_ms;          // current time
8     bool is_set_timer;      // whether start timer
9     TCPReceiverMessage recMsg; // the msg of TCPReceiver(ackno,window_size)
10    uint64_t abs_seqno;
11    size_t pre_window_size;
12    std::deque<TCPSenderMessage> outstanding_byte; // send but not ack
13    std::deque<TCPSenderMessage> ready_send_byte; // ready to send;
14    size_t outstanding_set_bytes;                // the bytes of outstanding
15    set
16    size_t retransmission_count;                // the count of
17    retransmission
18 public:
19     /* Construct TCP sender with given default Retransmission Timeout and
20     possible ISN */
21     TCPSender( uint64_t initial_RTO_ms, std::optional<Wrap32> fixed_isn );
22
23     /* Push bytes from the outbound stream */
24     void push( Reader& outbound_stream );
25
26     /* Send a TCPSenderMessage if needed (or empty optional otherwise) */
27     std::optional<TCPSenderMessage> maybe_send();
28
29     /* Generate an empty TCPSenderMessage */
30     TCPSenderMessage send_empty_message() const;
31
32     /* Receive an act on a TCPReceiverMessage from the peer's receiver */
33     void receive( const TCPReceiverMessage& msg );
34
35     /* Time has passed by the given # of milliseconds since the last time the
36     tick() method was called. */
37     void tick( uint64_t ms_since_last_tick );
38
39     /* Accessors for use in testing */
40     uint64_t sequence_numbers_in_flight() const; // How many sequence numbers
41     are outstanding?
```

```
38  uint64_t consecutive_retransmissions() const; // How many consecutive
    *re*transmissions have happened?
39  };
```

▼ tcp_sender.cc

```
1  /* TCPSender constructor (uses a random ISN if none given) */
2  TCPSender::TCPSender( uint64_t initial_RTO_ms, optional<Wrap32> fixed_isn )
3      : isn_( fixed_isn.value_or( Wrap32 { random_device()() } ) )
4      , is_send_ISN( false )
5      , is_send_FIN( false )
6      , initial_RTO_ms_( initial_RTO_ms )
7      , cur_RTO_ms( initial_RTO_ms )
8      , is_set_timer( false )
9      , recMsg()
10     , abs_seqno( 0 )
11     , pre_window_size( 1 )
12     , outstanding_byte()
13     , ready_send_byte()
14     , outstanding_set_bytes( 0 )
15     , retransmission_count( 0 )
16 {
17     recMsg.ackno = isn_;
18     recMsg.window_size = 1;
19 }
20
21 uint64_t TCPSender::sequence_numbers_in_flight() const
22 {
23     // Your code here.
24     return outstanding_set_bytes;
25 }
26
27 uint64_t TCPSender::consecutive_retransmissions() const
28 {
29     // Your code here.
30     return retransmission_count;
31 }
32
33 optional<TCPSenderMessage> TCPSender::maybe_send()
34 {
35     // Your code here.
36     if ( ready_send_byte.size() == 0 )
37         return nullopt;
38     TCPSenderMessage msg( ready_send_byte.front() );
39     ready_send_byte.pop_front();
40     is_set_timer = true;
41     return msg;
42 }
```

```

43
44 void TCPSender::push( Reader& outbound_stream )
45 {
46     // Your code here.
47     while ( outstanding_set_bytes < recMsg.window_size ) {
48         // read a msg
49         TCPSenderMessage msg;
50         // 1.whether have send?
51         if ( !is_send_ISN ) {
52             is_send_ISN = true;
53             msg.SYN = true;
54             msg.seqno = isn_;
55         } else {
56             msg.seqno = Wrap32::wrap( abs_seqno, isn_ );
57         }
58
59         // 2. set lenth
60         size_t data_len
61             = min( min( static_cast<size_t>( recMsg.window_size -
outstanding_set_bytes ), TCPConfig::MAX_PAYLOAD_SIZE ),
62                   static_cast<size_t>( outbound_stream.bytes_buffered() ) );
63         // 3. get data
64         read( outbound_stream, data_len, msg.payload );
65         // 4.
66         if ( outbound_stream.is_finished() && msg.sequence_length() +
outstanding_set_bytes < recMsg.window_size ) {
67             if ( !is_send_FIN ) {
68                 is_send_FIN = true;
69                 msg.FIN = true;
70             }
71         }
72         // 5.
73         if ( msg.sequence_length() == 0 )
74             break;
75         else {
76             outstanding_byte.push_back( msg );
77             ready_send_byte.push_back( msg );
78             outstanding_set_bytes += msg.sequence_length();
79         }
80         // 6.
81         abs_seqno += msg.sequence_length();
82
83         (void)outbound_stream;
84     }
85 }
86
87 TCPSenderMessage TCPSender::send_empty_message() const
88 {

```

```

89 // Your code here.
90 TCPSenderMessage msg;
91 msg.seqno = Wrap32::wrap( abs_seqno, isn_ );
92 return msg;
93 }
94
95 void TCPSender::receive( const TCPReceiverMessage& msg )
96 {
97 // Your code here.
98 recMsg = msg;
99 if ( recMsg.window_size == 0 )
100     recMsg.window_size = 1;
101 pre_window_size = msg.window_size;
102 if ( msg.ackno.has_value() ) {
103     if ( msg.ackno.value().unwrap( isn_, abs_seqno ) > abs_seqno )
104         return;
105
106     while ( outstanding_set_bytes != 0
107             && outstanding_byte.front().seqno.unwrap( isn_, abs_seqno ) +
outstanding_byte.front().sequence_length()
108             <= msg.ackno.value().unwrap( isn_, abs_seqno ) ) {
109         outstanding_set_bytes -= outstanding_byte.front().sequence_length();
110         outstanding_byte.pop_front();
111         if ( outstanding_set_bytes == 0 ) {
112             is_set_timer = false;
113         } else {
114             is_set_timer = true;
115         }
116         cur_RTO_ms = initial_RTO_ms_;
117         retransmission_count = 0;
118     }
119 }
120 (void)msg;
121 }
122
123 void TCPSender::tick( const size_t ms_since_last_tick )
124 {
125 // Your code here.
126 if ( is_set_timer ) {
127     cur_RTO_ms -= ms_since_last_tick;
128 }
129 if ( cur_RTO_ms <= 0 ) {
130     // retransmission
131     ready_send_byte.push_front( outstanding_byte.front() );
132     retransmission_count++;
133     // ARQ
134     if ( pre_window_size > 0 ) {
135         cur_RTO_ms = pow( 2, retransmission_count ) * initial_RTO_ms_;

```

```

136     } else {
137         cur_RTO_ms = initial_RTO_ms_;
138     }
139 }
140 (void)ms_since_last_tick;
141 }

```

▼

```

1 cmake --build build --target check3
2 Test project /home/sgt/cs/minnow/build
3     Start 1: compile with bug-checkers
4 1/36 Test #1: compile with bug-checkers ..... Passed    14.53 sec
5     Start 3: byte_stream_basics
6 2/36 Test #3: byte_stream_basics ..... Passed     0.01 sec
7     Start 4: byte_stream_capacity
8 3/36 Test #4: byte_stream_capacity ..... Passed     0.01 sec
9     Start 5: byte_stream_one_write
10 4/36 Test #5: byte_stream_one_write ..... Passed     0.02 sec
11     Start 6: byte_stream_two_writes
12 5/36 Test #6: byte_stream_two_writes ..... Passed     0.02 sec
13     Start 7: byte_stream_many_writes
14 6/36 Test #7: byte_stream_many_writes ..... Passed     0.07 sec
15     Start 8: byte_stream_stress_test
16 7/36 Test #8: byte_stream_stress_test ..... Passed     0.47 sec
17     Start 9: reassembler_single
18 8/36 Test #9: reassembler_single ..... Passed     0.02 sec
19     Start 10: reassembler_cap
20 9/36 Test #10: reassembler_cap ..... Passed     0.02 sec
21     Start 11: reassembler_seq
22 10/36 Test #11: reassembler_seq ..... Passed     0.03 sec
23     Start 12: reassembler_dup
24 11/36 Test #12: reassembler_dup ..... Passed     0.04 sec
25     Start 13: reassembler_holes
26 12/36 Test #13: reassembler_holes ..... Passed     0.02 sec
27     Start 14: reassembler_overlapping
28 13/36 Test #14: reassembler_overlapping ..... Passed     0.02 sec
29     Start 15: reassembler_win
30 14/36 Test #15: reassembler_win ..... Passed     5.48 sec
31     Start 16: wrapping_integers_cmp
32 15/36 Test #16: wrapping_integers_cmp ..... Passed     0.02 sec
33     Start 17: wrapping_integers_wrap
34 16/36 Test #17: wrapping_integers_wrap ..... Passed     0.01 sec
35     Start 18: wrapping_integers_unwrap
36 17/36 Test #18: wrapping_integers_unwrap ..... Passed     0.01 sec
37     Start 19: wrapping_integers_roundtrip

```

```

38 18/36 Test #19: wrapping_integers_roundtrip ..... Passed 1.40 sec
39      Start 20: wrapping_integers_extra
40 19/36 Test #20: wrapping_integers_extra ..... Passed 0.27 sec
41      Start 21: recv_connect
42 20/36 Test #21: recv_connect ..... Passed 0.02 sec
43      Start 22: recv_transmit
44 21/36 Test #22: recv_transmit ..... Passed 0.49 sec
45      Start 23: recv_window
46 22/36 Test #23: recv_window ..... Passed 0.02 sec
47      Start 24: recv_reorder
48 23/36 Test #24: recv_reorder ..... Passed 0.02 sec
49      Start 25: recv_reorder_more
50 24/36 Test #25: recv_reorder_more ..... Passed 10.74 sec
51      Start 26: recv_close
52 25/36 Test #26: recv_close ..... Passed 0.02 sec
53      Start 27: recv_special
54 26/36 Test #27: recv_special ..... Passed 0.03 sec
55      Start 28: send_connect
56 27/36 Test #28: send_connect ..... Passed 0.02 sec
57      Start 29: send_transmit
58 28/36 Test #29: send_transmit ..... Passed 0.59 sec
59      Start 30: send_retx
60 29/36 Test #30: send_retx ..... Passed 0.02 sec
61      Start 31: send_window
62 30/36 Test #31: send_window ..... Passed 0.31 sec
63      Start 32: send_ack
64 31/36 Test #32: send_ack ..... Passed 0.02 sec
65      Start 33: send_close
66 32/36 Test #33: send_close ..... Passed 0.02 sec
67      Start 34: send_extra
68 33/36 Test #34: send_extra ..... Passed 0.10 sec
69      Start 36: compile with optimization
70 34/36 Test #36: compile with optimization ..... Passed 3.06 sec
71      Start 37: byte_stream_speed_test
72          ByteStream throughput: 0.46 Gbit/s
73 35/36 Test #37: byte_stream_speed_test ..... Passed 0.35 sec
74      Start 38: reassembler_speed_test
75          Reassembler throughput: 0.40 Gbit/s
76 36/36 Test #38: reassembler_speed_test ..... Passed 0.61 sec
77
78 100% tests passed, 0 tests failed out of 36
79
80 Total Test time (real) = 38.90 sec
81 Built target check3

```


